

## OR 610 – Assignment 02

Abhishek Shambhu

### Code with commented outputs:

```

"""
OR 610 Assignment 02 Solution
@author: Abhishek Shambhu
"""

####Question:

#1. Organic DL Without using a deep learning framework, code a one-hidden layer neural
network model
#(multi-layer perceptron) for classification, i.e. the last layer is softmax and with
#cross-entropy loss function. Implement back-propagation. Use your code to classify MNIST.
#You can download MNIST data from http://yann.lecun.com/exdb/mnist/ and read it as a
numpy array
#using the following function.

#import struct import numpy as np
#def read_idx(filename):
#    with open(filename, 'rb') as f:
#        zero, data_type, dims = struct.unpack('>HBB', f.read(4))
#        shape = tuple(struct.unpack('>I', f.read(4))[0] for d in range(dims))
#        return np.fromstring(f.read(), dtype=np.uint8).reshape(shape)

#Run a few numerical experiments with different non-linear activation functions,
#different number of neurons and different learning rates.

#2. Regularization Add l2 regularization to your model. Compare the accuracy and training
speed.

####Solution:
#1.One hidden layer neural network without using a Deep Learning framework
import struct
import numpy as np
import time

def read_idx(filename):
    with open(filename, 'rb') as f:
        zero, data_type, dims = struct.unpack('>HBB', f.read(4))
        shape = tuple(struct.unpack('>I', f.read(4))[0] for d in range(dims))
        return np.frombuffer(f.read(), dtype=np.uint8).reshape(shape)

X_train = read_idx(r'D:\sem3\or610\HW2\train-images.idx3-ubyte')
y_train = read_idx(r'D:\sem3\or610\HW2\train-labels.idx1-ubyte')
X_test = read_idx(r'D:\sem3\or610\HW2\t10k-images.idx3-ubyte')
y_test = read_idx(r'D:\sem3\or610\HW2\t10k-labels.idx1-ubyte')

```

```

class MLP:
    """ Defining a One-Hidden Layer Neural Network Model (Multi-Layer Perceptron) for
    classification
    Parameters:
        n: number of iterations
        l_rate: learning rate
        hidden_neurons: number of hidden neurons
        act_func: activation function used on the hidden layer
        encode: One hot encode target values to binary form if not
        size: size of each batch (number of samples in each batch)
    """
    def __init__(self, n=1000, l_rate=0.05, hidden_neurons=10, act_func = 'sigmoid', encode=True,
size=1):
        self.n = n
        self.l_rate = l_rate
        self.encode=encode
        self.hidden_neurons = hidden_neurons
        self.size = size
        self.act_func = act_func

    def hidden_layer(self, z):
        if self.act_func=='relu':
            return np.clip(z, 0, 250)
        else:
            return 1/(1+ np.exp(-z))

    def softmax(self, z):
        z = z- np.max(z)
        return np.exp(z)/np.sum(np.exp(z), axis=1).reshape(-1,1)

    def onehot(self, data):
        n_classes = np.unique(data)
        onehot_array = np.zeros((data.shape[0], np.max(data)+1))
        onehot_array[np.arange(data.shape[0]), data] = 1
        if n_classes.shape[0] > 2:
            return onehot_array[:,n_classes]
        else:
            return onehot_array[:,n_classes[0]]

    def forward_prop(self, X):
        Zh = np.dot(X, self.Wh) + self.bias_h
        Ah = self.hidden_layer(Zh)
        Zout = np.dot(Ah, self.Wout) + self.bias_out
        Aout = self.softmax(Zout)
        return Zh, Ah, Zout, Aout

    def backward_prop(self, Zh, Ah, Zout, Aout, X_train, y_train):
        output_error = Aout - y_train
        DJ_DWout = np.dot(Ah.T, output_error)

```

```

if self.act_func == 'relu':
    hidden_deriv = Ah
    hidden_deriv[hidden_deriv > 0] = 1
    hidden_deriv[hidden_deriv <= 0] = 0
else:
    hidden_deriv = Ah * (1-Ah)
DJ_DWhid = np.dot(output_error, self.Wout.T) * hidden_deriv
grad_Whid = np.dot(X_train.T, DJ_DWhid)
grad_Bhid = np.sum(DJ_DWhid, axis=0)
grad_Wout = DJ_DWout
grad_Bout = np.sum(output_error, axis=0)
self.Wh -= grad_Whid * self.l_rate
self.bias_h -= grad_Bhid * self.l_rate
self.Wout -= grad_Wout * self.l_rate
self.bias_out -= grad_Bout * self.l_rate

def fit(self, X_train, y_train, X_valid, y_valid):
    if self.encode:
        y_train_oh = self.onehot(y_train)
        n_classes = y_train_oh.shape[1]
        n_features = X_train.shape[1]
        self.Wh = np.random.randn(n_features, self.hidden_neurons)
        self.Wout = np.random.randn(self.hidden_neurons, n_classes)
        self.bias_h = np.zeros((1, self.hidden_neurons))
        self.bias_out = np.zeros((1, n_classes))
        m = y_train.shape[0]
        self.metrics = {'cost': [], 'train_accuracy': [], 'valid_accuracy': []}

    for i in range(self.n):
        shuffled_values = np.random.permutation(m)
        X_shuffled = X_train[shuffled_values]
        y_shuffled = y_train_oh[shuffled_values]
        for batch in range(0, m, self.size):
            x_batch = X_shuffled[batch:batch+self.size]
            y_batch = y_shuffled[batch:batch+self.size]
            Zh, Ah, Zout, Aout = self.forward_prop(x_batch)
            self.backward_prop(Zh, Ah, Zout, Aout, x_batch, y_batch)

        Zh, Ah, Zout, Aout = self.forward_prop(X_train)
        cost = self.cost_function(Aout, y_train_oh)
        train_predictions = self.predict(Aout)

        Zh, Ah, Zout, Aout = self.forward_prop(X_valid)
        valid_predictions = self.predict(Aout)

        train_accuracy = np.sum(train_predictions ==
y_train).astype(np.float)/train_predictions.shape[0]
        valid_accuracy = np.sum(valid_predictions ==
y_valid).astype(np.float)/valid_predictions.shape[0]

```

```

        if not (i+1)%20:
            print("Iterations: {} \t Training Accuracy: {:.3f} \t Validation Accuracy: {:.3f}".format(i+1,
train_accuracy, valid_accuracy))
            self.metrics['cost'].append(cost)
            self.metrics['train_accuracy'].append(train_accuracy)
            self.metrics['valid_accuracy'].append(valid_accuracy)
        return self

def cost_function(self, Aout, y):
    return np.average(-y*np.log(Aout) - ((1-y)*np.log(1-Aout)))

def predict(self, output):
    return np.argmax(output, axis=1)

#2.One hidden layer neural network without using a Deep Learning framework and using L2
Regularization
class MLP_l2:
    """ Defining a One-Hidden Layer Neural Network Model (Multi-Layer Perceptron) for
classification with l2 regularization
Parameters:
    l2: l2 regularization rate
    n: number of iterations
    l_rate: learning rate
    hidden_neurons: number of hidden neurons
    act_func: activation function used on the hidden layer
    encode: One hot encode target values to binary form if not
    size: size of each batch (number of samples in each batch)
    """
    def __init__(self, l2=0.0, n=1000, l_rate=0.05, hidden_neurons=10, act_func = 'sigmoid',
encode=True, size=1):
        self.l2 = l2
        self.n = n
        self.l_rate = l_rate
        self.encode=encode
        self.hidden_neurons = hidden_neurons
        self.size = size
        self.act_func = act_func

    def hidden_layer(self, z):
        if self.act_func=='relu':
            return np.clip(z, 0, 250)
        else:
            return 1/(1+ np.exp(-z))

    def softmax(self, z):
        z = z- np.max(z)
        return np.exp(z)/np.sum(np.exp(z), axis=1).reshape(-1,1)

    def onehot(self, data):
        n_classes = np.unique(data)

```

```

onehot_array = np.zeros((data.shape[0], np.max(data)+1))
onehot_array[np.arange(data.shape[0]), data] = 1
if n_classes.shape[0] > 2:
    return onehot_array[:,n_classes]
else:
    return onehot_array[:,n_classes[0]]

def forward_prop(self, X):
    Zh = np.dot(X, self.Wh) + self.bias_h
    Ah = self.hidden_layer(Zh)
    Zout = np.dot(Ah, self.Wout) + self.bias_out
    Aout = self.softmax(Zout)
    return Zh, Ah, Zout, Aout

def backward_prop(self, Zh, Ah, Zout, Aout, X_train, y_train):
    output_error = Aout - y_train
    DJ_DWout = np.dot(Ah.T, output_error)
    if self.act_func == 'relu':
        hidden_deriv = Ah
        hidden_deriv[hidden_deriv > 0] = 1
        hidden_deriv[hidden_deriv <= 0] = 0
    else:
        hidden_deriv = Ah * (1-Ah)
    DJ_DWhid = np.dot(output_error, self.Wout.T) * hidden_deriv
    grad_Whid = np.dot(X_train.T, DJ_DWhid)
    grad_Bhid = np.sum(DJ_DWhid, axis=0)
    grad_Wout = DJ_DWout
    grad_Bout = np.sum(output_error, axis=0)
    l2_Whid = self.Wh * self.l2
    l2_Wout = self.Wout * self.l2
    self.Wh -= (grad_Whid + l2_Whid) * self.l_rate
    self.bias_h -= grad_Bhid * self.l_rate
    self.Wout -= (grad_Wout * self.l_rate) + l2_Wout
    self.bias_out -= grad_Bout * self.l_rate

def fit(self, X_train, y_train, X_valid, y_valid):
    if self.encode:
        y_train_oh = self.onehot(y_train)
        n_classes = y_train_oh.shape[1]
        n_features = X_train.shape[1]
        self.Wh = np.random.randn(n_features, self.hidden_neurons)
        self.Wout = np.random.randn(self.hidden_neurons, n_classes)
        self.bias_h = np.zeros((1,self.hidden_neurons))
        self.bias_out = np.zeros((1,n_classes))
        m = y_train.shape[0]
        self.metrics = {'cost': [], 'train_accuracy': [], 'valid_accuracy': []}

    for i in range(self.n):
        rearrange_values = np.random.permutation(m)
        X_shuffled = X_train[rearrange_values]

```

```

y_shuffled = y_train_oh[rearrange_values]
for batch in range(0, m, self.size):
    x_batch = X_shuffled[batch:batch+self.size]
    y_batch = y_shuffled[batch:batch+self.size]
    Zh, Ah, Zout, Aout = self.forward_prop(x_batch)
    self.backward_prop(Zh, Ah, Zout, Aout, x_batch, y_batch)
Zh, Ah, Zout, Aout = self.forward_prop(X_train)
cost = self.cost_function(Aout, y_train_oh)
train_predictions = self.predict(Aout)
Zh, Ah, Zout, Aout = self.forward_prop(X_valid)
valid_predictions = self.predict(Aout)
train_accuracy = np.sum(train_predictions ==
y_train).astype(np.float)/train_predictions.shape[0]
valid_accuracy = np.sum(valid_predictions ==
y_valid).astype(np.float)/valid_predictions.shape[0]

    if not (i+1)%20:
        print("Iterations: {} \t Training Accuracy: {:.3f} \t Validation Accuracy: {:.3f}".format(i+1,
train_accuracy, valid_accuracy))
        self.metrics['cost'].append(cost)
        self.metrics['train_accuracy'].append(train_accuracy)
        self.metrics['valid_accuracy'].append(valid_accuracy)
    return self

def cost_function(self, Aout, y):
    m = y.shape[0]
    l2_cost = (self.l2/(2*m))*(np.sum(self.Wh**2)+np.sum(self.Wout**2))
    return np.average(-y*np.log(Aout) - ((1-y)*np.log(1-Aout))) + l2_cost

def predict(self, output):
    return np.argmax(output, axis=1)

## Scaling to adjust weights
X_train = X_train.astype(np.float)/255.
X_test = X_test.astype(np.float)/255.

## Reshaping
X_train = X_train.reshape(-1,784)
X_test = X_test.reshape(-1,784)

# Running a NN without L2 Regularization and check accuracy and speed
#Sigmoid hidden activation, 100 hidden layers, 0.001 learning rate, 1000 batch size
model1 = MLP(n=100,hidden_neurons=100, size=1000, l_rate=0.001)
start = time.time()
model1.fit(X_train=X_train[:50000],
    y_train=y_train[:50000],
    X_valid=X_train[50000:],
    y_valid=y_train[50000:])
end =time.time()
runtime = end - start

```

```
print("Runtime for NN without L2 Regularization having Sigmoid hidden activation, 100 hidden layers, 0.001 learning rate, 1000 batch size:" + str(runtime) + "s")
```

```
'''
```

```
Iterations: 20 Training Accuracy: 0.914 Validation Accuracy: 0.917
Iterations: 40 Training Accuracy: 0.938 Validation Accuracy: 0.932
Iterations: 60 Training Accuracy: 0.950 Validation Accuracy: 0.939
Iterations: 80 Training Accuracy: 0.958 Validation Accuracy: 0.943
Iterations: 100 Training Accuracy: 0.964 Validation Accuracy: 0.946
Runtime for NN without L2 Regularization having Sigmoid hidden activation, 100 hidden layers, 0.001 learning rate, 1000 batch size:71.7282361984253s
```

```
'''
```

```
#ReLU Activation Function, 100 hidden layers, 0.001 learning rate, 1000 batch size
model2 = MLP(n=100,hidden_neurons=100, size=1000, l_rate=0.001, act_func='relu')
```

```
start = time.time()
```

```
model2.fit(X_train=X_train[:50000],
```

```
        y_train=y_train[:50000],
```

```
        X_valid=X_train[50000:],
```

```
        y_valid=y_train[50000:])
```

```
end =time.time()
```

```
runtime = end - start
```

```
print("Runtime for NN without L2 Regularization having ReLU Activation Function, 100 hidden layers, 0.001 learning rate, 1000 batch size:" + str(runtime) + "s")
```

```
'''
```

```
Iterations: 20 Training Accuracy: 0.906 Validation Accuracy: 0.904
Iterations: 40 Training Accuracy: 0.931 Validation Accuracy: 0.921
Iterations: 60 Training Accuracy: 0.947 Validation Accuracy: 0.936
Iterations: 80 Training Accuracy: 0.954 Validation Accuracy: 0.942
Iterations: 100 Training Accuracy: 0.959 Validation Accuracy: 0.943
Runtime for NN without L2 Regularization having ReLU Activation Function, 100 hidden layers, 0.001 learning rate, 1000 batch size:66.75671648979187s
```

```
'''
```

```
#Sigmoid hidden activation, 30 hidden layers, 0.001 learning rate, 1000 batch size
```

```
model3 = MLP(n=100,hidden_neurons=30, size=1000, l_rate=0.001)
```

```
start = time.time()
```

```
model3.fit(X_train=X_train[:50000],
```

```
        y_train=y_train[:50000],
```

```
        X_valid=X_train[50000:],
```

```
        y_valid=y_train[50000:])
```

```
end =time.time()
```

```
runtime = end - start
```

```
print("Runtime for NN without L2 Regularization having Sigmoid hidden activation, 30 hidden layers, 0.001 learning rate, 1000 batch size:" + str(runtime) + "s")
```

```
'''
```

```
Iterations: 20 Training Accuracy: 0.888 Validation Accuracy: 0.895
Iterations: 40 Training Accuracy: 0.913 Validation Accuracy: 0.913
```

```

Iterations: 60 Training Accuracy: 0.925 Validation Accuracy: 0.921
Iterations: 80 Training Accuracy: 0.934 Validation Accuracy: 0.926
Iterations: 100 Training Accuracy: 0.939 Validation Accuracy: 0.931
Runtime for NN without L2 Regularization having Sigmoid hidden activation, 30 hidden layers,
0.001 learning rate, 1000 batch size:47.72391986846924s
'''

```

```

#ReLU Activation Function, 30 hidden layers, 0.001 learning rate, 1000 batch size
model4 = MLP(n=100,hidden_neurons=30, size=1000, l_rate=0.001, act_func='relu')
start = time.time()
model4.fit(X_train=X_train[:50000],
          y_train=y_train[:50000],
          X_valid=X_train[50000:],
          y_valid=y_train[50000:])
end = time.time()
runtime = end - start
print("Runtime for NN without L2 Regularization having ReLU Activation Function, 30 hidden
layers, 0.001 learning rate, 1000 batch size:" + str(runtime) + "s")
'''

```

```

Iterations: 20 Training Accuracy: 0.867 Validation Accuracy: 0.876
Iterations: 40 Training Accuracy: 0.896 Validation Accuracy: 0.901
Iterations: 60 Training Accuracy: 0.915 Validation Accuracy: 0.920
Iterations: 80 Training Accuracy: 0.917 Validation Accuracy: 0.922
Iterations: 100 Training Accuracy: 0.927 Validation Accuracy: 0.929
Runtime for NN without L2 Regularization having ReLU Activation Function, 30 hidden layers,
0.001 learning rate, 1000 batch size:48.041693449020386s
'''

```

```

#Sigmoid hidden activation, 300 hidden layers, 0.0001 learning rate, 1000 batch size
model5 = MLP(n=100,hidden_neurons=300, size=1000, l_rate=0.0001)
start = time.time()
model3.fit(X_train=X_train[:50000],
          y_train=y_train[:50000],
          X_valid=X_train[50000:],
          y_valid=y_train[50000:])
end = time.time()
runtime = end - start
print("Runtime for NN without L2 Regularization having Sigmoid hidden activation, 300 hidden
layers, 0.0001 learning rate, 1000 batch size:" + str(runtime) + "s")
'''

```

```

Iterations: 20 Training Accuracy: 0.886 Validation Accuracy: 0.891
Iterations: 40 Training Accuracy: 0.913 Validation Accuracy: 0.913
Iterations: 60 Training Accuracy: 0.925 Validation Accuracy: 0.923
Iterations: 80 Training Accuracy: 0.933 Validation Accuracy: 0.928
Iterations: 100 Training Accuracy: 0.938 Validation Accuracy: 0.933
Runtime for NN without L2 Regularization having Sigmoid hidden activation, 300 hidden layers,
0.0001 learning rate, 1000 batch size:42.71775031089783s
'''

```



```
#ReLU Activation Function, 300 hidden layers, 0.0001 learning rate, 1000 batch size
model6 = MLP(n=100,hidden_neurons=300, size=1000, l_rate=0.0001, act_func='relu')
start = time.time()
model4.fit(X_train=X_train[:50000],
          y_train=y_train[:50000],
          X_valid=X_train[50000:],
          y_valid=y_train[50000:])
end =time.time()
runtime = end - start
print("Runtime for NN without L2 Regularization having ReLU Activation Function, 300 hidden
layers, 0.0001 learning rate, 1000 batch size:" + str(runtime) +"s")
```

```
'''
```

```
Iterations: 20  Training Accuracy: 0.881    Validation Accuracy: 0.882
Iterations: 40  Training Accuracy: 0.917    Validation Accuracy: 0.920
Iterations: 60  Training Accuracy: 0.929    Validation Accuracy: 0.929
Iterations: 80  Training Accuracy: 0.935    Validation Accuracy: 0.934
Iterations: 100 Training Accuracy: 0.940    Validation Accuracy: 0.935
Runtime for NN without L2 Regularization having ReLU Activation Function, 300 hidden layers,
0.0001 learning rate, 1000 batch size:40.78611636161804s
```

```
'''
```

```
# Running a NN with L2 Regularization and check accuracy and speed
#Sigmoid hidden activation, 30 hidden layers, 0.001 learning rate, 300 batch size, 0.1 l2
model7 = MLP_l2(n=100,hidden_neurons=30, size=300, l_rate=0.001, l2=0.1)
start = time.time()
model7.fit(X_train=X_train[:50000],
          y_train=y_train[:50000],
          X_valid=X_train[50000:],
          y_valid=y_train[50000:])
end =time.time()
runtime = end - start
print("Runtime for NN with L2 Regularization having Sigmoid hidden activation, 30 hidden
layers, 0.001 learning rate, 300 batch size:" + str(runtime) +"s")
```

```
'''
```

```
Iterations: 20  Training Accuracy: 0.615    Validation Accuracy: 0.625
Iterations: 40  Training Accuracy: 0.688    Validation Accuracy: 0.701
Iterations: 60  Training Accuracy: 0.752    Validation Accuracy: 0.768
Iterations: 80  Training Accuracy: 0.700    Validation Accuracy: 0.716
Iterations: 100 Training Accuracy: 0.762    Validation Accuracy: 0.776
Runtime for NN with L2 Regularization having Sigmoid hidden activation, 30 hidden layers,
0.001 learning rate, 300 batch size:42.96507692337036s
```

```
'''
```

```
#Sigmoid hidden activation, 30 hidden layers, 0.001 learning rate, 300 batch size, 0.5 l2
model8 = MLP_l2(n=100,hidden_neurons=30, size=300, l_rate=0.001, l2=0.5)
start = time.time()
model8.fit(X_train=X_train[:50000],
```

```
y_train=y_train[:50000],
X_valid=X_train[50000:],
y_valid=y_train[50000:])
end =time.time()
runtime = end - start
print("Runtime for NN with L2 Regularization having Sigmoid hidden activation, 30 hidden
layers, 0.001 learning rate, 300 batch size:" + str(runtime) +"s")
```

```
'''
```

```
Iterations: 20  Training Accuracy: 0.344    Validation Accuracy: 0.346
Iterations: 40  Training Accuracy: 0.291    Validation Accuracy: 0.283
Iterations: 60  Training Accuracy: 0.441    Validation Accuracy: 0.449
Iterations: 80  Training Accuracy: 0.496    Validation Accuracy: 0.512
Iterations: 100 Training Accuracy: 0.374    Validation Accuracy: 0.379
Runtime for NN with L2 Regularization having Sigmoid hidden activation, 30 hidden layers,
0.001 learning rate, 300 batch size:43.55163931846619s
'''
```

```
'''
```

#### Findings:

1. With lower L2 regularization rates model accuracy is better.
2. Sigmoid and ReLU Activation function seems to have almost the same effect on the model.

```
'''
```

**Summary of results:**

<b>Model</b>	<b>Parameters</b>	<b>Training Accuracy</b>	<b>Validation Accuracy</b>	<b>Time</b>
Model1	Sigmoid hidden activation, 100 hidden layers, 0.001 learning rate, 1000 batch size	0.964	0.946	71.7282361984253s
Model2	ReLU Activation Function, 100 hidden layers, 0.001 learning rate, 1000 batch size	0.959	0.943	66.75671648979187s
Model3	Sigmoid hidden activation, 30 hidden layers, 0.001 learning rate, 1000 batch size	0.939	0.931	47.72391986846924s
Model4	ReLU Activation Function, 30 hidden layers, 0.001 learning rate, 1000 batch size	0.927	0.929	48.041693449020386s
Model5	Sigmoid hidden activation, 300 hidden layers, 0.0001 learning rate, 1000 batch size	0.938	0.933	42.71775031089783s
Model6	ReLU Activation Function, 300 hidden layers, 0.0001 learning rate, 1000 batch size	0.940	0.935	40.78611636161804s
Model7	Sigmoid hidden activation, 30 hidden layers, 0.001 learning rate, 300 batch size, 0.1 l2	0.762	0.776	42.96507692337036s
Model8	Sigmoid hidden activation, 30 hidden layers, 0.001 learning rate, 300 batch size, 0.5 l2	0.374	0.379	43.55163931846619s